

SpectralBL Boundary Layer Analysis Report:

Reconstructing the Nocturnal Attractor Manifold

Automated Analytics Pipeline Engine
SpectralBL-Analytics Framework

June 25, 2026

Abstract

This report details the low-rank manifold reconstruction and structural stability profiling executed for the **ALL** campaign. Characterized by intense near-surface radiative cooling under strongly stable stratification, this boundary layer configuration represents a regime where classical surface-layer similarity theory frequently breaks down.

Using singular value decomposition (SVD) on multi-level vertical profile sequences, we extract an empirical low-rank state-space manifold that captures the dominant modes of the nocturnal boundary layer. Structural identifier optimization resolves the underlying vector fields, which are subsequently evaluated using Stage 5 parameter continuation.

By projecting high-dimensional vertical profile measurements directly onto empirical spatial basis vectors extracted via singular value decomposition (SVD), we reconstruct the underlying low-dimensional attractor manifold in the space (η_1, η_2, η_3) . This spatial-projection framework eliminates the need for delay-embedding preprocessing, instead capturing structural modes of the boundary layer column with direct physical interpretability. The result is a low-rank, continuous phase space for diagnosing boundary layer complexity and mapping regime transitions where traditional surface-flux similarity theories break down.

1 Introduction and Physical Context

Traditional boundary layer parameterizations rely heavily on local gradient profiles and standard Monin-Obukhov Similarity Theory (MOST). However, in the very stable boundary layer (vSBL), the coupling between the surface layer and the upper air column is frequently severed by the development of a strong thermal inversion. The atmosphere enters a regime dominated by radiative cooling, low-level jet (LLJ) development, global intermittency, and submesoscale internal gravity waves.

In these conditions, local turbulence metrics such as Ri_g become unreliable diagnostics: once Ri_g exceeds the critical stability threshold of 0.25, classical surface-layer similarity theory predicts the collapse of continuous turbulence. In practice, however, the boundary layer does not simply switch off. Intermittent bursting, wave-mode excitation, and the residual decay of turbulent kinetic energy continue in ways that local gradient metrics are fundamentally blind to.

To track these non-local structural reconfigurations, the SpectralBL pipeline extracts continuous spatial pattern modes of the boundary layer state. The resulting low-rank coordinates (η_1, η_2, η_3) correspond directly to deep physical phenomena: bulk background layer evolution, shear-driven jet modulation, and vertical wave-like structural curvature. These coordinates remain dynamically active throughout stable and super-critical- Ri_g conditions, providing information about boundary layer behavior that gradient-based scalings cannot access.

2 Observational Transitions and Regime Geometries

3 Transition Event Mapping

Transition plots are the primary dynamical fingerprint for regime transitions in the stable boundary layer. This section provides a three-panel bridge between continuation stability signatures and reduced-state trajectory structure for campaign ALL.

Transition assets were not available at render time. Expected files: ../../data/outputs/transition_panel_a_all.csv, ../../data/outputs/transition_panel_b_all.csv, and ../../data/outputs/transition_panel_c_all.csv.

4 Attractor Dynamics and Low-Dimensional Phase Space

4.1 Overview

The boundary layer is analyzed as a low-dimensional dynamical system through projection onto dominant singular vectors. The following section presents the reconstructed phase space trajectory for campaign `ALL` using the v0.1 baseline.

Campaign: `ALL`

Baseline Version: v0.1

Total Samples: 5600

Mean Entropy: $H_{\text{mean}} = 0.122$

4.2 3D Phase Space Projection

Figure 1 visualizes the three-dimensional trajectory (η_1, η_2, η_3) . To highlight the true multi-regime orbital evolution, timesteps are connected sequentially by a continuous historical path line. The trajectory is colored using a high-contrast colormap driven by the **Attractor Spin** diagnostic $(\Omega(t))$, explicitly isolating the directionality of regime transitions.




Figure 1: 3D attractor state orbit trajectory. Color variation tracks Attractor Spin (Ω) , separating shear breakout acceleration (cool blue) from radiative collapse configurations (warm red).

4.3 Physical Coordinate Interpretation

This phase-space view acts as the compressed state engine of the atmospheric boundary layer (ABL), projecting high-dimensional tower profiles into three dominant, orthogonal coordinates:

- η_1 (**Bulk Background Mode**): Captures the mean inversion strength, nocturnal cooling depth, and overall boundary-layer depth evolution.
- η_2 (**Shear & LLJ Mode**): Captures shear production, alongside low-level jet (LLJ) speed and height modulation.
- η_3 (**Vertical Structural Curvature Mode**): Captures submesoscale gravity-wave excitation and the precursor structural signatures of intermittent turbulence events.

4.3.1 Geometric Reading Guide

- **Tight closed loops** indicate near-repeatable diurnal cycling (convective daytime mixing to stable nighttime recovery and back).
- **Broad, tangled trajectories** indicate non-stationary forcing, intermittency, or mesoscale disruption.
- **Color evolution by $\Omega(t)$** isolates transition directionality: positive spin values (warm tones) indicate stable collapse paths; negative values (cool tones) mark turbulent regeneration breakouts.

4.4 Trajectory Metrics

The mean low-rank coefficients across all samples are:

$$\langle \eta_1 \rangle = -11.240 \tag{1}$$

$$\langle \eta_2 \rangle = -0.067 \tag{2}$$

$$\langle \eta_3 \rangle = 0.129 \tag{3}$$

The magnitude of the attractor vector $\|\boldsymbol{\eta}\|$ evolves continuously throughout the campaign, reflecting the transient structural adjustments driven by diurnal heating and nocturnal radiative cooling cycles.

Operationally, peaks in $\|\boldsymbol{\eta}\|$ often coincide with rapid boundary-layer reconfiguration, while spin excursions identify windows where similarity-based single-regime assumptions are least reliable.

4.5 Information-Theoretic Manifold Complexity Evaluation

To preserve algorithmic invariance across divergent energy scales, rank truncation is enforced with a machine-precision floor criterion ($s_i \geq \epsilon_{\text{mach}} s_1$). This execution yields:

- **Constrained modes:** $N_c = 3$
- **Unconstrained nullspace modes:** $N_0 = 0$
- **Condition number:** $\kappa = 22.52$
- **Shannon effective dimension:** $D_{\text{eff}} = 1.6010$

The entropy-derived $D_{\text{eff}} = 1.6010$ quantifies how many modes are actively carrying structure in practice, rather than only counting nominal basis size. Larger values indicate broader modal participation and lower compressibility of the observed boundary-layer state.

5 Spectral Geometry and Regime Structural Analysis

5.1 Mathematical Foundation

The decomposition of the boundary layer into regimes relies on the spectral properties of the low-rank attractor subspace. We define three primary diagnostic metrics:

1. **Singular Value Entropy** $H = -\sum_i p_i \log(p_i)$ where $p_i = \sigma_i^2 / \sum_j \sigma_j^2$
2. **Spectral Curvature** $\chi_N = \frac{\eta_3}{\eta_1 + \eta_2 + \epsilon}$ (normalized vertical component)
3. **Attractor Magnitude** $\|\boldsymbol{\eta}\| = \sqrt{\eta_1^2 + \eta_2^2 + \eta_3^2}$

5.2 Campaign Context

- **Campaign:** ALL
- **Baseline:** v0.1 (spectralbl-attractor)
- **Temporal Coverage:** 1463.5 hours
- **Sample Count:** 5600 observations

5.3 Entropy Evolution

The mean singular value entropy for this campaign is:

$$H_{\text{mean}} = 0.122$$

This value reflects the typical balance between coherent, low-dimensional structure and multi-scale turbulent activity. Nocturnal periods show reduced entropy (stratified regime), while daytime transitions display elevated entropy (mixing regime).

5.4 Structural Independence Verification

To verify that the low-rank decomposition captures physically independent aspects of boundary-layer dynamics, we contrast spectral geometry against classical stability diagnostics.

The gradient Richardson number can be interpreted as

$$\text{Ri}_g \sim \frac{\text{buoyant suppression}}{\text{shear production}}.$$

In practice, low Ri_g indicates active shear-driven turbulence, while large Ri_g indicates strongly stable conditions where continuous turbulence is difficult to sustain.

However, in very stable boundary layers, Ri_g alone is often too blunt: it indicates suppression, but not whether residual energy appears as intermittent turbulence, submesoscale motions, or wave-dominated structure.

The spectral curvature metric χ_N addresses this by tracking geometric structure in the reduced attractor space rather than relying only on local gradients.

When Ri_g exceeds the critical threshold (0.2–0.25), classical surface-layer similarity theory predicts that continuous turbulence should collapse entirely. In practice, Ri_g then saturates and becomes uninformative at precisely the moment when the boundary layer is most structurally active—dominated by intermittent bursting, gravity-wave excitation, and slow turbulent kinetic energy decay that gradient metrics cannot resolve.

The scatter plot shown in Figure 2 demonstrates this directly. While Ri_g stalls at high values, the vertical coordinate η_3 remains dynamically resolved throughout the same stability range, mapping the active wave field and intermittent bursting events that Ri_g misses entirely.

Scientific interpretation of orthogonality:

- If χ_N were redundant with Ri_g , points would collapse toward a single diagonal trend.
- Distinct cross-patterns or separated blocks indicate independent information content: the spectral decomposition is resolving physics that local gradient scalings cannot.

- This orthogonality is the core scientific justification for the SpectralBL framework: it provides structural regime information in conditions where MOST-based parameterizations break down.




Figure 2: Projected attractor coordinates with orthogonality structure visualized via η_3 color encoding.

5.5 Non-Orthogonal Scale Coupling Accounts

Because smooth partition-of-unity masks (ψ_M, ψ_W, ψ_T) retain localized overlap to stabilize regime transitions, total energy tracking is not strictly additive at every instant.

For this execution, the off-diagonal interaction proxy is

$$\mathcal{I}_{ij} = 0.000000.$$

This term is tracked to quantify cross-regime leakage (wave-turbulence exchange) during active stratification shifts and to preserve auditability of coupling behavior in the reduced-order representation.

Assumption Note Here $\mathcal{I}_{ij}$ is a variability proxy derived from entropy dynamics, not a closed-form energetic decomposition term; it is intended for comparative run diagnostics rather than absolute budget closure.

6 Geometric Precursor Analysis

6.1 Trajectory Geometry as Early-Warning Signal

We introduce trajectory curvature $\kappa(t) = \|\boldsymbol{\eta}' \times \boldsymbol{\eta}''\| / \|\boldsymbol{\eta}'\|^3$ as a coordinate-independent early-warning signal for regime transitions. This geometric

measure captures how sharply the phase-space trajectory turns in (η_1, η_2, η_3) coordinates, independent of the basis choice.

6.1.1 Multi-Scale Indicator Categories

Category A: Critical-Transition Indicators (classical theory validation)

- Multi-scale variance: $\text{Var}_\tau[\eta_3]$ for $\tau \in \{5, 15, 30, 60\}$ minutes
- Lag-1 autocorrelation: $\text{AC}_1^\tau(\eta_3)$ tests for critical slowing down

Category B: Geometric Indicators (novel in SBL context)

- Trajectory speed: $v(t) = \|\text{d}\boldsymbol{\eta}/\text{d}t\|$
- Acceleration magnitude: $a(t) = \|\text{d}^2\boldsymbol{\eta}/\text{d}t^2\|$
- Trajectory curvature: $\kappa(t)$ (regularized where $v > \epsilon$)

6.2 ALL Precursor Diagnostics

Campaign characteristics:

- Duration: 1463.5 hours (60.98 days)
- Samples: 5600 profiles at 15.7-minute intervals
- State-space range: $\eta_3 \in [-3.16, 2.45]$
- Mean trajectory speed: $\bar{v} = n/a$

6.2.1 Key Observations

6.3 Physical Interpretation

High curvature $\kappa \gg \bar{\kappa}$ indicates trajectory *sharply turning* in manifold space—not merely moving fast (speed v) or changing rapidly (acceleration a), but undergoing rapid rotation in the direction of motion. Physically, this corresponds to:

- Boundary layer restructuring events
- Navigation of folds in equilibrium manifold
- Transitions between distinct dynamical regimes
- Possible precursor to bifurcation (approaching/departing critical manifold regions)




Figure 3: Multi-panel geometric precursor time series for ALL. Panels show (top to bottom): Category A variance indicators, Category A autocorrelation, Category B speed & acceleration, Category B trajectory curvature (coordinate-independent), and classical Richardson number. Vertical markers indicate identified transition events.

7 Attractor Stability and Bifurcation Analysis

8 System Diagnostics and Definitive Conclusions

8.1 Pipeline Execution Verification

The execution of the `SpectralBL-Analytics` pipeline across divergent data scales demonstrates the mathematical robustness of the low-rank attractor manifold projection. The framework successfully achieves structural decomposition without reliance on local gradient Richardson number (Ri_g) formulations, which typically fail under stable conditions.

8.2 Comparative Information-Theoretic Summary

Based on the processed analytics run, the structural characteristics of the underlying dataset are summarized below:

Diagnostic Metric	Pipeline Evaluation Summary for ALL
Data Footprint	Analyzed 5600 observations across a temporal scale of 1463.5 hours.
Manifold Complexity	Mean singular value entropy registers at $H_{\text{mean}} = 0.122$, yielding a Shannon effective dimension of $D_{\text{eff}} = 1.6010$.
Numerical Stability	The condition number $\kappa = 22.52$ confirms a well-conditioned low-rank basis projection ($N_0 = 0$ nullspace modes).

Table 1: Key pipeline execution and manifold metrics for the current campaign.

8.3 Definitive Scientific Conclusions

By mapping the boundary layer dynamics into the orthogonal coordinates (η_1, η_2, η_3) , this analysis establishes the following definitive conclusions based on the campaign-specific footprint:

- **Composite Regime Envelope:** The merged trajectory captures a broader attractor envelope spanning process-oriented and climatological campaign characteristics.
- **Cross-Campaign Orthogonality:** Shared eta-space coordinates preserve campaign-specific signatures while maintaining a common comparison basis.
- **Validation Breadth:** Multi-campaign aggregation expands stress-testing of closures and parameterizations across contrasting stability environments.

8.4 Operational Framework Recommendation

These signatures confirm that `ALL` contributes campaign-specific constraints in the verification pipeline. Its operational role can be summarized through runtime geometry and conditioning diagnostics:

$$\mathcal{M}_{\text{validation}} = \mathbf{f}\left(\mathcal{A}_{\text{site}}(\kappa_{\text{rank}} = 22.52, D_{\text{eff}} = 1.6010), \mathcal{S}_{\text{surface}}\right) \quad (4)$$

Surface roughness context for this run: `0.050 m`.

When integrated into multi-campaign workflows, low-roughness and canopy-dominated sites jointly provide empirical constraints for model behavior in non-local, intermittent, and collapse-prone stable boundary-layer conditions.

9 Physical Interpretation of the Continuation Parameter

The continuation parameter γ represents the effective atmosphere–surface coupling strength inferred from the sparse structural operator identified in Stage 4. Small values of γ correspond to weak turbulent transport and strongly decoupled nocturnal conditions, where the surface inversion evolves in relative isolation from the overlying residual layer. Larger values of γ represent enhanced vertical exchange between the surface inversion and the residual layer above, sustaining coherent turbulent structures through the column.

Parameter continuation therefore traces the geometric deformation of the reconstructed attractor manifold as the atmospheric state transitions between coupled and decoupled regimes. The leading eigenvalue $\alpha(\gamma) = \max \Re(\lambda_i)$ serves as a scalar indicator of manifold stability: negative values indicate a contracting, stable attractor; the zero crossing at γ_c marks the onset of structural instability; and the imaginary part $\beta(\gamma) = |\text{Im}(\lambda_{\text{dom}})|$ tracks the intrinsic oscillatory frequency of the emerging mode.

10 Stage 5 Bifurcation Spectrum

Stage 5 branch artifacts were not available for this campaign at report render time. Expected source: `../data/outputs/stage5_bifurcation_branches_all.csv`.

Cross-campaign Stage 5 comparison table omitted because one or more required summary manifests were not available at render time.

A Pipeline Metadata and Verification Criteria

- Source code: <https://github.com/davidengland/SpectralBL-Analytics>
- Build timestamp: June 25, 2026
- Campaign: **ALL**
- Mathematical basis: p-FEM spectral projection and scale-invariant rank truncation
- Stability threshold for transition diagnostics: $Ri_c = 0.25$
- p-FEM basis/geometry dimensions are recorded in campaign stage diagnostics and operator build logs.
- Institutional Provenance: Core boundary-layer telemetry and analytical hosting courtesy of the University of Alabama in Huntsville (UAH), managed via the National Space Science and Technology Center (<https://nsstc.uah.edu>).
- Templating: Mustache – active
- Figure cache: TikZ external compilation – **tikz-cache/**